

BSc Project
**Learning to Play Tetris using the Covariance
Matrix Adaptation Evolution Strategy**

January 6, 2016

Contents

1	Abstract	3
2	Nomenclature	3
2.1	Abbreviations	3
2.2	Symbols	3
3	Introduction	4
3.1	Goals of the thesis	4
4	Previous work	5
5	Scope and limitations	6
6	Objective function	6
7	Optimizers	7
8	CE (Cross Entropy)	8
8.1	Input	9
8.2	Loop	10
9	CMA-ES	10
10	Experiments	10
10.1	MDP-Tetris	10
10.2	Normalization of samples	11
10.3	Setup	12
10.4	Assesment of controller performance	13
10.5	Verification of CE	14
10.6	Optimal settings for Cross Entropy	15
10.7	Comparison between CMA-ES and CE	16
10.8	Results	16
10.9	Analysis and discussion	16

1 Abstract

ABSTRACT

2 Nomenclature

2.1 Abbreviations

CMA-ES, Covariance Matrix Adaption Evolution Strategy

CE, Cross-Entropy

2.2 Symbols

m , the mean of a distribution.

M , the matrix used as second argument to the gaussian distribution, which describes the shape of the distribution. Usually, when a sample from a gaussian distribution is drawn, it's written as: $x \sim \mathcal{N}(\mu, \sigma^2)$, however in (Hansen, 2011), μ and σ is used for other purposes, and therefore these symbols are reserved for this.

N , number of agents generated per generation/iteration.

γ , number of games played by each agent for evaluation.

μ , number of the population samples included in the next generation.

O , the function that estimates the performance of an agent/sample point.

Z , noise term added to the variance, specific to Cross-Entropy.

t , counter indicating the current generation/iteration.

w , the weight configuration of a Tetris playing agent(n -dimensional vector).

F , a feature of the Tetris game state. This is a function that maps a feature of the Tetris board to a real value.

V , a linear combination of feature functions and weight. The function used for evaluating a Tetris game state. The function is composed of a linear combination of feature sets and their associated weights, accepts a game state, and yields a single real-value.

s , a state of the Tetris game board.

x , an individual vector from a generation/iteration. Represented with an n -dimensional vector.

3 Introduction

On the topic of reinforcement learning, a widely used benchmark for learning algorithms are designing agents for playing the classical game of Tetris. Tetris is an appealing benchmarking problem due to its complexity. The standard game plays on a board made from a grid that is 10 cells wide and 20 cells tall. As the game progresses, differently shaped pieces fall from the top of the board. When a row on the board is fully occupied by pieces, the line is removed, all lines above it moved one line down and a score point is given to the player. If a cell above the 20 rows first is occupied, the game ends. The task of the player is to move and rotate the falling pieces in a way that yields the highest score before the game ends.

Tetris is indeed a hard task to computationally optimize, as the game has a very high number of board configurations estimated to be 10^{59} (Scherrer, 2009b). Because of this complexity, a common approach in the literature is to use *one-piece controllers*¹, such as described in Scherrer (2009a). These controllers are aware of only the current board state and the currently falling piece. Using these controllers, the search space is reduced to only looking at the current board, and the possible places to drop the piece. The game used for the benchmark is a simplified version of Tetris, in which controllers need only to decide in what column to drop the current piece, and what orientation the piece should have when dropped. Thus, the simplified version of Tetris differs from the original game mainly in two significant aspects. First, the controller is disallowed to move the piece horizontally while the piece is falling. Second, the controller has 'infinite' time to make its decision on where to drop the current piece. Thus, the controller cannot take advantage of moving the piece during the fall, but is not restricted by the time limitations. This is however a common way of benchmarking Tetris controllers (Scherrer, 2009b)

When the controllers decide which action to take, it will simulate each of the possible actions and choose the one that leads to the most favourable board state. To evaluate the board state, the controller uses a set of features that defines various qualities of the board, and associate a weight to each feature. This means that the efficiency of the controller is determined by the features the controller is aware of and how heavily they are weighted. This allows the controller with n features to be expressed as an n dimensional real-valued vector, with one dimension per feature, and the value in that dimension the weight. An often referred to controller is the Dellacherie's controller, as described in Scherrer (2009b). This controller takes six features of the board into account, seen in table 1 on page 5.

3.1 Goals of the thesis

Both the Cross-Entropy and the CMA-ES methods have been used in learning Tetris with *one piece controllers*, but as mentioned, to our knowledge, only little effort has been put into comparing the two methods. This thesis will explore how the two methods compare against each other under similar conditions. Therefore, we will use a set of features among those commonly used, and compare how the two optimizing algorithms differ. The goal however is not to find a controller that outperforms existing controllers, but only to investigate how the Cross-Entropy and CMA-ES differs when learning Tetris with similar controllers.

¹Agents and controllers both refer to artificial players.

In this paper, we will explore how the two state-of-the-art optimization algorithms, Cross-Entropy and CMA-ES, differ when applied to the task of playing Tetris.

The SHARK library (Igel et al., 2008) contains a working implementation of the CMA-ES algorithm. However, the Cross-Entropy method is not present in the SHARK library and thus, a part of the thesis is to implement it ourselves according to the other researchers work. To document the soundness of the implemented CE method, we will replicate the experiment in (Thiery and Scherrer, 2009) and verify that we obtain the similar results.

Then, we will benchmark CMA-ES and CE against each other to determine if one yields better optimization results than the other.

4 Previous work

Over the time, numerous researchers has tried different feature sets and applied various optimizers to find the best possible Tetris controllers. The features used are typically ones that attempt to mimic the board conditions that would normally catch the attention of a human player, such as how high the overall pile of pieces is and how many holes the board has. In (Scherrer, 2009a) table 1, a table presents some feature sets used throughout various publications on the subject. In later works, many authors have had success with applying evolutionary stochastic search methods for tuning the weights of the feature sets towards efficient controllers. For the purpose of this thesis, we are in particular addressing the Cross-Entropy method described in detail in (Pieter-Tjerk De Boer and Rubenstein, 2014) and the Covariance Matrix Adaption Evolution Strategy (CMA-ES), described in (Hansen, 2011). The particular Cross-Entropy method applied is the one described in (Szita and Lörincz, 2006) as the "Noisy Cross Entropy Method".

Feature	Description
Landing height	The height of the piece when it lands
Eroded piece cells	Number of rows cleared in the last move times the number of bricks cleared from the last move
Row transitions	Number of horizontal cell transitions
Column transitions	Number of vertical cell transitions
Holes	Number of empty cells covered by a full cell
Board wells	Cumulative sum of cells to the depth of the board wells.

Figure 1: features of the Dellacherie controller

Currently, many researchers have proposed numerous feature sets and multiple optimization methods have been explored. A controller often referred to is the Dellacherie controller (Fahy, 2003). This controller was hand-tuned by trial and error, and did originally, on a regular non-simplified Tetris game, achieve an average of 660 000 lines. The same feature set (see figure 1) is often incorporated in later works when optimizing controllers. An earlier feature set is the set proposed by (Bertsekas and Tsitsiklis, 1996) referred to as Bertsekas and Tsitsiklis features. In 2006, Szita and Lörincz (Szita and Lörincz, 2006) applied the Cross-

Entropy method using the Bertsekas and Tsitsiklis features. They report that using no noise, their controller converged at 300 000 lines on average. The best result reported in (Szita and Lörincz, 2006) is when decreasing noise is applied, in which the controller’s score exceeded 800 000 lines. However, in a later paper, using Dellacherie, Bertsekas and two selfdefined features achieved 35.000.000 lines $\pm 20\%$ (Scherrer, 2009b).

Creating a Tetris-controller is a NP-complete problem, where we want to find a strategy which maximizes the average score. Most researchers utilize three general approaches to create policies.

- Handwritten controllers
- Reinforcement learning approaches
- Optimization algorithms

But as seen in (Scherrer, 2009b), a combination of the methods can yield good results, where Thiery and Scherrer employed Dellacherie’s handwritten policy and used CE to optimize it.

5 Scope and limitations

The experiments will be carried out on the simplified version of Tetris using the MDP_Tetris software found at (MDP_Tetris, 2009), which is the same simulator used by Scherrer et al. in (Scherrer, 2009a) among other authors in various papers. This software already have the well known feature sets implemented, so we will not ourselves extend any of the features. The source code of the Tetris simulator is used as-is, and is therefore not altered prior to running the experiments. For comparing the optimizers, the SHARK (Igel et al., 2008) library will be used. This library already contains an implementation of the CMA-ES optimizer, but lacks the Cross-Entropy. Therefore, a part of this thesis will be to implement and document the Cross-Entropy method in Shark.

6 Objective function

The optimizing algorithms used in this thesis both attempts to optimize a certain function. In this case, the optimization function aims to develop the best possible Tetris playing agent. Thus, the value of the objective function describes an estimated performance of the input agent. The objective function serves as an abstraction to the Tetris emulator that will evaluate the agent by letting the agent play a pre-configured number of games. The objective function that estimates the performance of an agent is later referred to as O .

When the Tetris simulator plays Tetris, the internal decision process of the Tetris controller is configured with a set of parameters which remain fixed across a single game. These parameters are the weights associated with each feature function that is considered by the controller. The features F_i each map a state s of the game to a real value. An overview of the exact mappings can be seen in table 1 in (Scherrer, 2009a). How much each of these mappings should affect the final evaluation of the board is determined by w_i denoting the weight of the

i -th feature. Finally, the function to assess the value of current board state $V(s)$, with n features functions present, can be expressed as:

$$V(s) = \sum_{i=1}^n w_i F_i(s)$$

Let S be the set of all states that each possible action can lead to from the current state s . The controller will then for each reachable state $gameState_i \in S$ evaluate the state by $O(s_i)$. The chosen action is the one that yields the state of the highest value. The performance of the controller is hence directly tied to the features and weights in the evaluation function. To adjust these controllers, one can either change the set of applied feature functions, or as the optimization algorithms will do, change the weighting of the features. For the experiments, the feature sets remain fixed, and the task of the optimization algorithm applied is to tune the set of weights in order to maximize the performance of the controller.

In summary, the objective function accepts a vector of values for each weight and configures an agent with the weights corresponding to the entries in the input vector. It then plays γ games with the agent and reports the mean score.

7 Optimizers

In this thesis, the Cross-Entropy method and the Covariance Matrix Evolution Strategy are compared. Both of the methods fall into the category of *stochastic optimization* methods. These methods are useful for optimization problems that have no gradient. The optimization functions aim to optimize the parameter set \mathbf{x} for the objective function O .

$$\hat{\mathbf{x}} = \arg \max_{\mathbf{x}} O(\mathbf{x}) \quad O : \mathbb{R}^n \rightarrow \mathbb{R}$$

In these methods, the optimizing algorithm uses a family of parametric distributions, and maintain a mean m along with other parameters to search the best possible solution for the objective function. In the case studied in this thesis both the CMA-ES and Cross-entropy methods use a Gaussian distribution to sample solutions to the objective function. Hence, both of the functions aim to find a mean m and an $n \times n$ matrix M^2 , such that when a vector \mathbf{x} is sampled by $\mathbf{x} \sim \mathcal{N}(m, M)$, then $O(\mathbf{x})$ is likely to yield preferable results.

The algorithms work iteratively, such that the mean and variance of the distribution is altered for each iteration t . The algorithms start by initializing the parameters either at random or some fixed point. A common configuration is setting the mean to all zeros and the standard deviation to the identity matrix. Thus, for the first iteration $t = 0$, a configuration could be:

$$m_0 = \begin{bmatrix} 0 \\ \vdots \\ 0 \end{bmatrix}, \quad M_0 = \begin{bmatrix} 1 & \dots & 0 \\ \vdots & \ddots & \vdots \\ 0 & \dots & 1 \end{bmatrix}$$

Where the subscript notes that the values occur in iteration 0.

²In (Hansen, 2011), σ is used for step-size in CMA-ES, so M is instead introduced as an arbitrary $n \times n$ matrix in its place.

In each iteration, the algorithms sample λ vectors and evaluate their fitness against the objective function. When each of the solutions are evaluated, they are ordered according to their fitness:

$$\{\mathbf{x}_1, \dots, \mathbf{x}_N\} \text{ Such that } f(\mathbf{x}_1) \geq f(\mathbf{x}_2), \dots, f(\mathbf{x}_{N-1}) \geq f(\mathbf{x}_N)$$

The mean and standard deviation for the next iteration, that is m_{t+1} and M_{t+1} is then updated usually by considering the best of the ordered solutions. How exactly these parameters are updated is individual for each method and can be seen in the following sections.

8 CE (Cross Entropy)

CE is described through many papers in slightly different ways. Using similar format as in the paper (Thiery and Scherrer, 2009).

This method uses a Gaussian distribution and attempts to find distribution parameters that yields good candidates for the objective function O .

The Cross-Entropy method starts with an initial mean m and standard deviation M . The mean is usually an n dimensional vector set to:

$$m = \begin{pmatrix} 0 \\ \vdots \\ 0 \end{pmatrix}$$

The variance is kept individual for each dimension, and is usually initialized as follows:

$$M = \begin{bmatrix} \sigma_1 & \dots & 0 \\ \vdots & \ddots & \vdots \\ 0 & \dots & \sigma_n \end{bmatrix}$$

Where in this context, σ represents *standard deviation*.

The algorithm then works iteratively on generations of individual search points acting as candidate inputs for the objective function. In each generation, N vectors are sampled by $x_i \sim \mathcal{N}(m, M^2)$, $i \in \{1, \dots, N\}$. The vectors are all evaluated against the fitness function and ordered such that $O(x_1) \geq \dots \geq O(x_N)$, and the μ best are chosen for updating the distribution parameters. The mean is updated as the centroid of the chosen vectors, and the variance is updated as the variance of the chosen vector in each dimension.

The pseudo code and details of the algorithm can be seen in figure 2 on page 9.

input

O : The function that estimates the performance of a vector x

(m_0, M_0^2) : The mean and variance of the initial distribution

N : The number of vectors sampled per generation/iteration

μ : The number of offspring selected for the new mean

Z_t : The noise added to each generation/iteration

loop

Generate N vectors x_1, x_2, \dots, x_N from $\mathcal{N}(m_t, M_t^2)$

Evaluate each vector using O

Select the μ vectors with the highest evaluation

Update m_{t+1} of the μ best vectors

Update M_{t+1}^2 of the μ best vectors + Z_t

end loop

Figure 2: The pseudo code for the Cross-Entropy algorithm

8.1 Input

The objective function

The function used to assess the value of a sampled vector. As described in the 'Optimizers' section, CE is a general stochastic iterative algorithm that tries to solve an optimization problem of the form (Thiery and Scherrer, 2009):

$$\hat{\mathbf{x}} = \arg \max_{\mathbf{x}} O(\mathbf{x})$$

Where x corresponds to a given vector, and O is our actual objective function.

The mean and variance of the gaussian distribution

Here m_t is the mean and M_t^2 is the variance of the gaussian distribution (m_t, M_t^2) . More specifically this gaussian distribution is defined as

$$\mathcal{N}(m_t, M_t^2)$$

Where t denotes the current iteration.

The number of vectors

N is the number of vectors sampled in each generation.

The number of offspring

μ is the number of vectors which are used to compute the new mean, m_{t+1} , and variance, M_{t+1}^2 , for next generation/iteration. These offspring vectors gets selected directly by taking μ vectors which got the best evaluation.

The noise factor

The noise factor, Z_t , is the amount of noise which is applied to the variance M^2 in iteration/-generation t . In general, noise is used to avoid the risk of a local optimum. There are different

kinds of noise settings, such as: no noise, constant noise and linear decreasing noise (Szita and Lörincz, 2006). When using no noise, Z_t is simply set to zero. When using constant noise, the same value is added to the variance M^2 in each iteration/generation. When using linear decreasing noise, Z_t is defined as $Z_t = \max(5 - t/10, 0)$.

8.2 Loop

Sampling the population

The first step of the loop is to create the new generation consisting of N vectors. These vectors are sampled randomly within the distribution $x_i \sim \mathcal{N}(m_t, M_t^2)$.

Evaluating the population

After sampling the population, the algorithm needs to order the vectors to find the μ best vectors, each vector x_i , $i \in \{1, \dots, N\}$ is evaluated using O . The value from the objective function then yields the estimated performance of each individual.

Selecting the offspring

As each x_i has an assigned evaluation value, and the μ best vectors gets selected by taking the x_i vectors with the highest evaluation value.

Updating the distribution parameters

When updating the distribution parameters for the next iteration (m_{t+1}, M_{t+1}^2) , the mean is updated by computing the centroid of the μ best vectors. This is formally defined as:

$$m_{t+1} := \frac{\sum_i^\mu x_i}{\mu}$$

The variance M_{t+1}^2 is updated to match the variance of the μ best vectors, such that the variance in dimension i matches the variance of the μ in dimension i . This is formally defined as:

$$M_{t+1}^2 := \frac{\sum_i^\mu (x_i - m_{t+1})^T (x_i - m_{t+1})}{\mu} + Z_{t+1}$$

9 CMA-ES

10 Experiments

This section will walk through our experimental setup, our verification of CE and and a analysis/discussion of the results.

10.1 MDP-Tetris

When running the experiments, the source code of the MDP-Tetris (MDPTetris, 2009) was used to emulate the Tetris games. The source code is accompanied with files that describe

the various existing features. These files contains the identifiers of each feature to use, as well as two numbers respectively describing the agents reward function and how to evaluate a game over state. The number for the reward function has remained unchanged at 0 during all experiments. The "game-over" evaluation was for the Bertsekas feature set initially set to 0. Setting the "game-over" evaluation to 0 means that the agent will not distinguish between regular moves and moves that results in losing the game. When running the experiments with this setting, a large portion of the agents never exceeded a zero mean score. However, setting the value to -1 , meaning that a "game-over" move yields $-\infty$ reward, none of the experiments got stuck on only zero scores. An example of the layout of the feature file can be seen in figure 3.

```

0      <- Describes the reward function
-1     <- Actions leading to game over is avoided at all cost
22     <- The policy contains 22 features
8 0    <- The feature with id 8 initially has weight 0
...    <- The remaining 21 features

```

Figure 3: Example of a file that describes a feature set.

10.2 Normalization of samples

As mentioned by some authors (Boumaza, 2009), the vector that describes the agent can very well be normalized such that the vector is a point that lies on the n -dimensional hypersphere.

The reason for this lies in the nature of the evaluation function. When the controller chooses an action, it will evaluate all the possible actions possible with the current piece. It will use the value function V of each state s_i and choose the state with the highest value from the value function. Thus, if the states are ordered such that:

$$V(s_1) > \dots > V(s_N)$$

The agent then chooses the action that transitions from the current state to state s_1 . Since the value function assess the state by the following:

$$V(s) = \sum_{i=1}^n w_i F_i(s)$$

Then scaling the input of the agent, the weight vector w , by a number $a \in \mathbb{R}$, $a > 0$ the assessment is changed by:

$$\begin{aligned} \sum_{i=1}^n a w_i F_i(s) &= a \sum_{i=1}^n w_i F_i(s) \\ &= a V(s) \end{aligned}$$

And the ordering remains:

$$a V(s_1) > \dots > a V(s_N)$$

Thus the order of the value functions of each state does not change, and the same s_1 is still chosen for any $a \in \mathbb{R}$, $a > 0$.

To verify this, the Tetris objective function was executed with the same vector and the same seed for the random generator with a scale $a \in \{0.1, 0.2, 0.3, \dots, 9.8, 9.9, 10.0\}$, and the agent scored exactly the same for each scale.

This can be used in experiments for various reasons. As reported in (Boumazza, 2009), normalizing the samples will prevent CMA-ES from diverging in step size, and it can prevent loosing precision if the magnitude of weight vector becomes larger than feasible for the used floating point number and avoids size limitations.

10.3 Setup

When executing the experiments, various parameters each have an impact on the final result of the learning curve. Thus, the parameters are adjusted, first to match the experiments run by other researchers, and later to conduct as fair as possible comparisons between Cross-Entropy and CMA-ES.

The amount of vectors sampled in each generation N has a high impact on the algorithm performance. By setting N high, more policies are evaluated per iteration, and leads to a more thorough exploration of the search space. Thus the higher N increases the chances of finding a better mean for the next iteration. However, higher N also results in the need for more evaluations per iteration. The goal for tuning this parameter is then to set N high enough to ensure exploration of good solutions, and yet low enough to avoid unnecessary evaluations.

In the implementation of CMA-ES from SHARK (Igel et al., 2008) , the algorithm itself determines the value of N according to the size of the search space. Cross-Entropy however, does not seem to have a general rule for this parameter, so this value is manually adjusted to fit the problem as well as possible.

As both of the optimizing algorithm uses a subset of the sampled vectors from a generation to update the distribution parameters, the number of offspring μ influences how the next generation is sampled. By setting the value too high, the algorithm risks ceasing to progress any further since the updated mean would be too close to the previous one to significantly make a difference. By setting the value too low, the risk of reaching a local optimum increases since the high-scoring agents might have reached their high performance by chance.

The CMA-ES itself manages setting μ and Cross-Entropy is set according to the problem. Most authors that uses Cross-Entropy for Tetris sets the offspring size to 10% of population size, that is $\mu = \lfloor 0.1 \cdot N \rfloor$.

The number of games, γ , is the number of games which each agent plays in each iteration. An agent's score is defined as the mean of the score of these γ games. We want this value low as possible, because as with the number of agents, N , The number of games, γ , is another major factor in the run-time of the algorithm. As Tetris is stochastic by nature, the score deviates a lot, even when the same agent with the same policy plays multiple games. Hence, when assessing the true performance of a policy it's rarely enough to play just few

games. Thus, setting γ high increases the likelihood of correctly choosing the best agents, yet, it also causes longer run times of the experiments.

Specific to the Cross-Entropy method, most authors report that the performance of the algorithm increases dramatically when the sampling distribution is associated with a noise term. The different types of noise are described in section 8. The noise term is adjusted in order to prevent the algorithm from reaching a local optimum. The current research shows that noise terms of $Z_t = 4$ and $Z_t = \max(5 - t/10)$ (Thiery and Scherrer, 2009) produces the best results. The constant noise (such as $Z_t = 4$) ensures that the algorithm never settles in a too small area from which it samples, and forces it to explore solutions that are further away from the mean. The further the algorithm progresses, the less noise is assumed needed, as the mean should approach a global optimum. to address this, the linear decreasing noise is applied as it will lower the noise term as the algorithm progresses.

For the various experiments, these parameters will be tuned for the specific purpose at hand. In the verification of the Cross-Entropy, the parameters are set to match those reported in similar papers (Thiery and Scherrer (2009), Szita and Lörincz (2006)). In the comparison of the two algorithms, the parameters will be set such that the Cross-Entropy operates under as similar conditions as CMA-ES, to ensure an unbiased comparison.

10.4 Assesment of controller performance

The performance of a one-piece controller has a very high variation, and is in other research verified to be exponentially distributed. As a result of the high variance of the controllers, the performance of single controllers are often presented along with a confidence interval for the estimated mean score of the controller. The estimated mean of the controllers score is calculated by:

$$\hat{m} = \sum_{i=1}^N x_i$$

Thus, the maximum likelihood estimation of the rate parameter³ of the distribution is given by

$$\hat{\lambda} = \frac{1}{\hat{m}}$$

A confidence interval is found by the following:

$$\frac{2N}{\hat{\lambda}\chi_{1-\frac{\alpha}{2}, 2N}^2} < \frac{1}{\lambda} < \frac{2N}{\hat{\lambda}\chi_{\frac{\alpha}{2}, 2N}^2}$$

However, for these experiments, an approximation for a 95% confidence on lower and upper bound of the rate parameter λ is used:

³Note that λ is in this context not, the population size but instead the rate parameter for the exponential distribution.

$$\lambda_{low} = \hat{\lambda} \left(1 - \frac{1.96}{\sqrt{N}} \right)$$

$$\lambda_{upp} = \hat{\lambda} \left(1 + \frac{1.96}{\sqrt{N}} \right)$$

By this, the 95% confidence interval for the mean m is:

$$\frac{1}{\lambda_{low}} < m < \frac{1}{\lambda_{upp}}$$

When a controllers score is presented as " $s \pm p$ " this means has an empirical mean score of s and a real mean that is with 95% likelihood within $s \pm p$.

ADD REFERENCE TO EXPONENTIAL DISTRIBUTIONS

10.5 Verification of CE

Because the SHARK (Igel et al., 2008) library already contains an implementation of CMA-ES, but not an implementation of CE, we extended the library with our own implementation of the algorithm. This ensures as many similar conditions as possible for the two optimization algorithms as possible.

In order to verify the correctness of the implementation, we used the same experiments as used by Christophe Thiery and Bruno Scherrer (Thiery and Scherrer, 2009). These experiments were used by Thiery and Scherrer to verify their own CE implementation with various types of noise correction. Therefore, we will perform the same experiments to verify our own contribution to the SHARK (Igel et al., 2008) library, by trying to achieve the same results.

The setup is mirrored from the paper (Thiery and Scherrer, 2009), with 100 agents ($N = 100$) per iteration. The 10 best agents ($\mu = 10$) will be used to update gaussian distribution. After each iteration, an agent with the updated mean plays 30 games and the mean of these scores are recorded for the learning curve.

During evaluation each agent plays one game, that is $\gamma = 1$.

Figure 5, 6 and 7 shows 10 runs of each noise type. Figure 4 shows the mean graph for each of the noise types. The goal of these experiments were to replicate the experiments reported in (Thiery and Scherrer, 2009). As the results seen from our experiments to a high degree resemble those reported by Thiery et. al, we conclude that our Cross-Entropy implementation works similar to theirs.

When evaluating the score of the agent we also want to compute the confidence interval in verifying the implementation of CE. Since the population size is static for the mean agent, the confidence also becomes static for all experiments with 30 games, which is $\pm 36\%$. Compared to other papers, this confidence interval is similar (Scherrer, 2009b).

Based on the mean graphs and confidence interval compared to other papers, we can hereby verify that our implementation of CE works as intended.

SPECIFIC EXAMPLES OF MEAN SCORE + CONFIDENCE INTERVAL




Figure 4: Cross-Entropy mean performance

10.6 Optimal settings for Cross Entropy

Other researchers run the Cross Entropy algorithm with population size of $N = 100$ and an offspring corresponding to 10% of the population size, resulting in $\mu = 10$. As it's not discussed why this exact setting is applied, various settings of the Cross-Entropy was run to asses the performance of other configurations. The experiments includes different population sizes $N \in \{10, 22, 50, 100, 200\}$ and offspring sizes of either 10% and 50%. A summary of the experiments can be seen in figure 8 on page 15.

N	μ	mean	Q1	Q2	Q3
-----	-------	------	----	----	----

Figure 8: Cross Entropy configuration test

The experiments with different population and offspring sizes does not seem to support a choice for any other configuration than the mostly commonly used $N = 100$ and $\mu = 10$.

With a configuration of $N = 50$ and $\mu = 5$ convergence is achieved faster. This means that the score limit is reached faster, which results in longer computation time, than the $N = 100$ and $\mu = 10$ configuration.

ADD GRAPHS OF THE TWO DIFFERENT CONFIGURATIONS




Figure 5: No noise

10.7 Comparison between CMA-ES and CE

ADD EXPERIMENTS THAT COMPARE CMA-ES TO CROSS ENTROPY

10.8 Results

SHOW RESULTS FROM COMPARISON EXPERIMENTS

10.9 Analysis and discussion

DISCUSS THE RESULTS OF THE COMPARISON EXPERIMENTS

11 Conclusion

CAN WE CONCLUDE IF EITHER ALGORITHM PERFORMS BEST?




Figure 6: Constant noise




Figure 7: Linear decreasing noise

References

- Bertsekas, D. and Tsitsiklis, J. (1996). Neuro-dynamic programming. *Athena Scientific*.
- Boumaza, A. (2009). On the evolution of artificial tetris players. *HAL*, 14:1–14.
- Fahey, C. (2003). Tetris ai. <http://www.colinfahey.com/tetris/>. Accessed: 2015-11-21.
- Hansen, N. (2011). The CMA evolution strategy: A tutorial.
- Igel, C., Heidrich-Meisner, V., and Glasmachers, T. (2008). Shark. *Journal of Machine Learning Research*, 9:993–996.
- MDPTetris (2009). mdptetris. *None*.
- Pieter-Tjerk De Boer, Dirk P. Kroese, S. M. and Rubenstein, R. Y. (2014). A tutorial on the cross-entropy method. *Springer Science*, 521:20–67.
- Scherrer, C. T. B. (2009a). Building controllers for tetris. *International Computer Games Association Journal*, 32:3–11.
- Scherrer, C. T. B. (2009b). Improvements on learning tetris with cross entropy. *International Computer Games Association Journal*, 32:23–33.
- Szita, I. and Lörincz, A. (2006). Learning tetris using the noisy cross-entropy method. *Neural Computation*, 18(12):2936–2941.
- Thiery, C. and Scherrer, B. (2009). Improvements on learning tetris with cross entropy. *International Computer Games Association Journal*, 32(1):23–33.

□